Introduction to

Algorithm Design and Analysis

[7] Selection

Yu Huang

http://cs.nju.edu.cn/yuhuang
Institute of Computer Software
Nanjing University

In the last class…

- **MergeSort**
  o Design
  o Cost – time & space

- **Lower bounds for** *comparison-based sorting*
  o Worst-case
  o Average-case

The Selection

- **Selection – warm-ups**
  - Finding *max* and *min*
  - Finding the *second largest* key
- **Adversary argument and lower bound**

- **Selection – select the *median***
  - Expected linear time
  - Worst-case linear time
- **A Lower Bound for Finding the Median**

The Selection Problem

- **Problem definition**
  - Suppose E is an array containing n elements with keys from some linearly order set, and let k be an integer such that $1 \le k \le n$. The selection problem is to find an element with the k^{th} smallest key in E.

- **Special cases**
  - Find the max/min – $k=n$ or $k=1$
  - Find the *median* $(k = \frac{n}{2})$

Selection vs. Searching
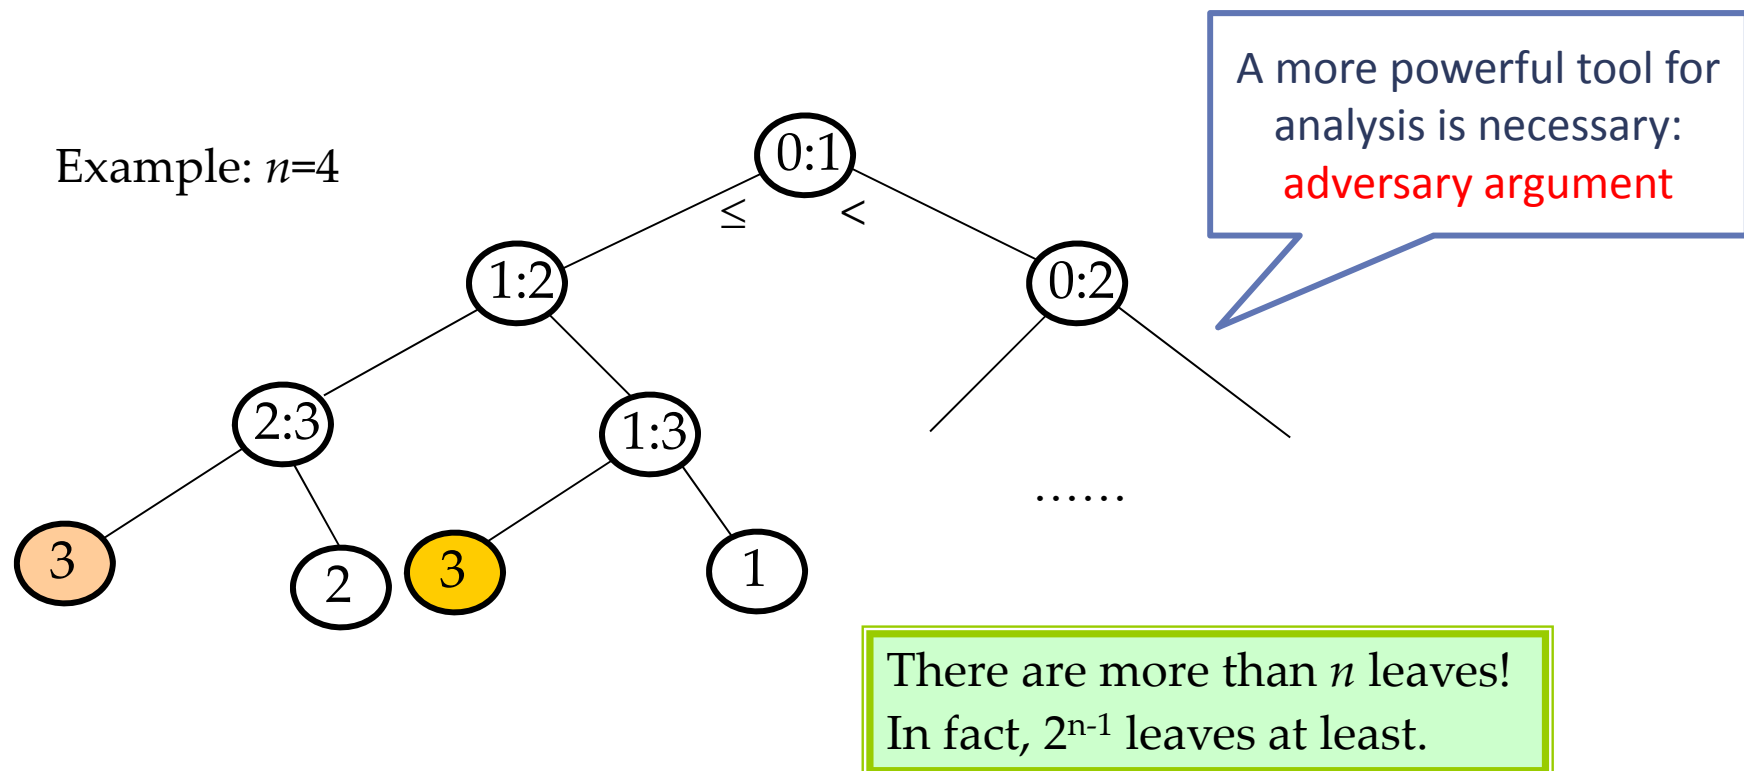
Lower Bound of Finding the Max

- **For any algorithm $\mathcal{A}$ that can compare and copy numbers exclusively, in the worst case, $\mathcal{A}$ cannot do fewer than $n\text{-}1$ comparisons to find the largest entry in an array with n entries.**

  - Proof: an array with n distinct entries is assumed. We can exclude a specific entry from being the largest entry only after it is determined to be "loser" to at least one entry. So, n-1 entries must be "losers" in comparisons done by the algorithm. However, each comparison has only one loser, so at least n-1 comparisons must be done.

Decision Tree and Lower Bound

Since the decision tree for the selection problem must have at least n leaves, the height of the tree is at least $\lceil \log n \rceil$. **It's not a good lower bound.**

Example: $n=4$

A more powerful tool for analysis is necessary: adversary argument

There are more than n leaves! In fact, 2^{n-1} leaves at least.

Finding *max* and *min*

- **The strategy**

  o Pair up the keys, and do $n/2$ comparisons(if n odd, having $E[n]$ uncompared);

  o Doing findMax for larger key set and findMin for small key set respectively (if n odd, $E[n]$ included in both sets)

- **Number of comparisons**

  o For even n: $n/2+2(n/2-1)=3n/2-2$

  o For odd n: $(n-1)/2+2((n-1)/2+1-1)=\lceil 3n/2 \rceil-2$

  How to prove this lower bound?   **Adversary Argument !**

Unit of Information

- **Max and Min**

  o That x is *max* can only be known when it is sure that every key other than x has lost some comparison.

  o That y is *min* can only be known when it is sure that every key other than y has win some comparison.

- **Each win or loss is counted as one unit of information**

  o *Any* algorithm must have at least 2n-2 units of information to be sure of specifying the *max* and *min*.

Adversary Strategy

Status of keys x and y Compared by an algorithm	Adversary response	New status	Units of new information
N,N	$x>y$	W,L	2
W,N or WL,N	$x>y$	W,L or WL,L	1
L,N	$x<y$	L,W	1
W,W	$x>y$	W,WL	1
L,L	$x>y$	WL,L	1
W,L or WL,L or W,WL	$x>y$	No change	0
WL,WL	Consistent with Assigned values	No change	0

The principle: let the key win if it never lose, or,
let the key lose if it never win, and

change one value if necessary

Lower Bound by the Adversary Argument

- **Construct an input to force *the* algorithm to do more comparisons as possible**
  - To give away as few as possible units of new information with each comparison.
    - It can be achieved that 2 units of new information are given away only when the status is N,N.
    - It is *always* possible to give adversary response for other status so that at most one new unit of information is given away, *without any inconsistencies*.
- **So, the *Lower Bound* is *n*/2+n-2(for even *n*)**

$$\frac{n}{2} \times 2 + (n - 2) \times 1 = 2n - 2$$

An Example

Comparison	x_1		x_2		x_3		x_4		x_5		x_6	
	S	V	S	V	S	V	S	V	S	V	S	V
x_1,x_2			N	*	N	*	N	*	N	*	N	*
x_1,x_5									L	5		
x_3,x_4												
x_3,x_6												12
x_3,x_1												
x_2,x_4												
x_5,x_6										WL	5	L	3
x_6,x_4								L	2			WL	3

Now, x_3 is the only one which never loses, so, Max is x_3

8 comparisons!
The lower bound is 7.

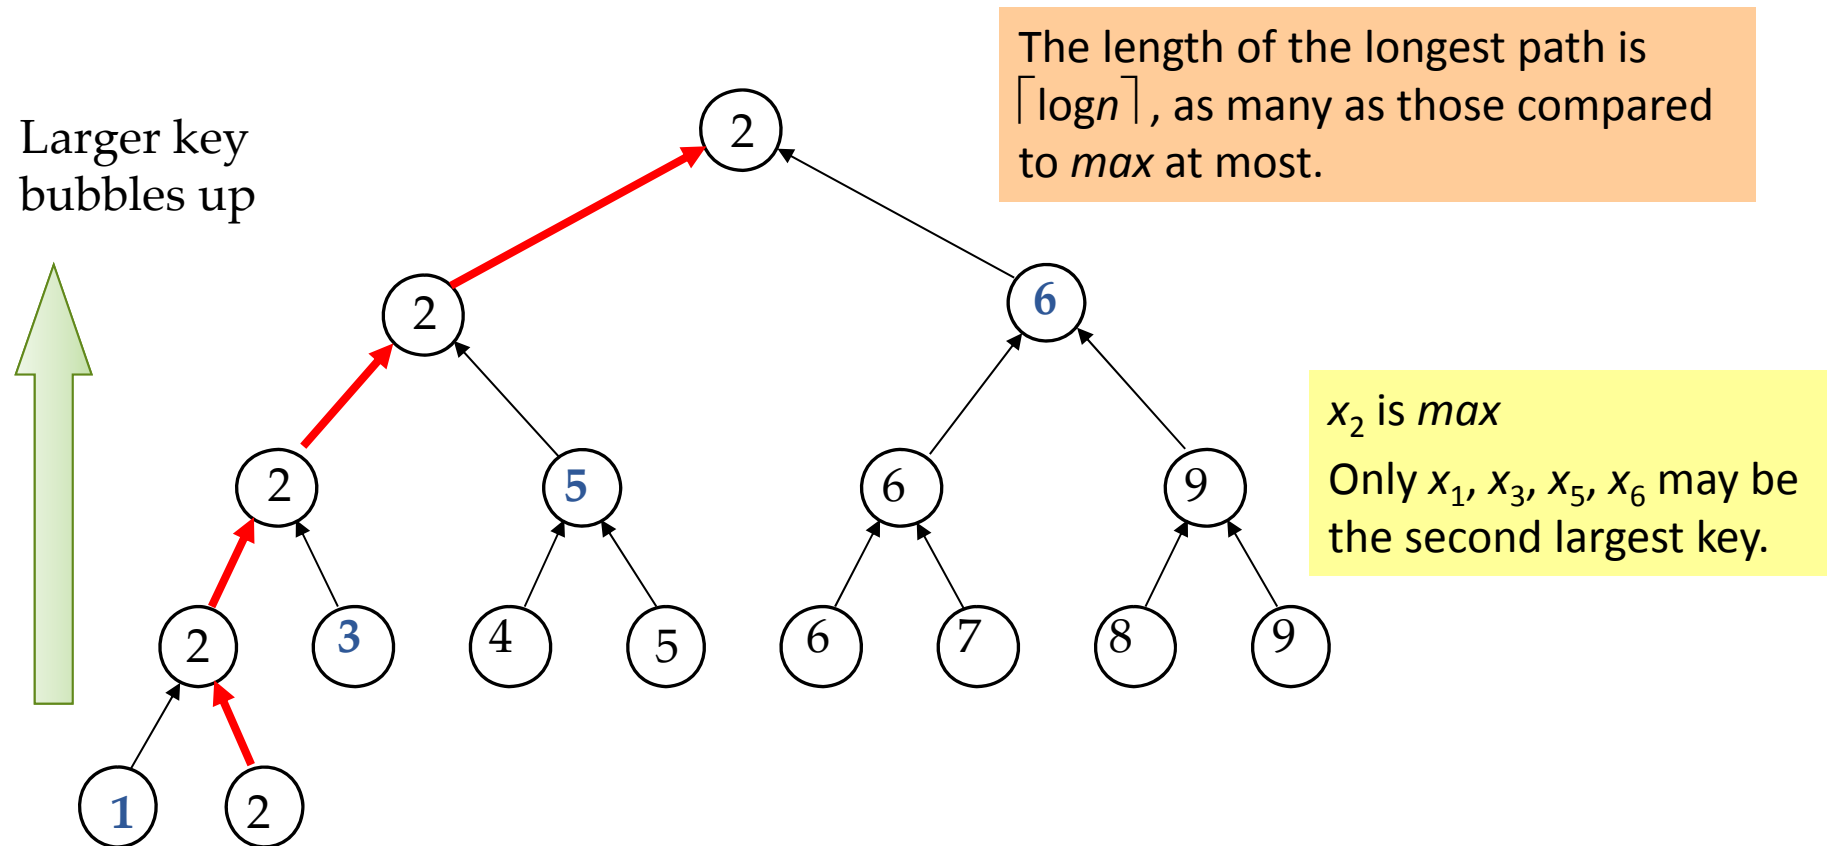Raising/lowering the value according to strategy

Find the 2nd Largest Key

- **Brute force - using FindMax twice**
  - Need 2n-3 comparisons.

- **For a better algorithm**
  - Collect some useful information from the first FindMax

- **Observations**
  - The key which loses to a key other than max cannot be the 2nd largest key.
  - To check "whether you lose to max?"

Tournament for the 2ⁿᵈ Largest Key

Larger key bubbles up

The length of the longest path is $\lceil \log n \rceil$, as many as those compared to *max* at most.

x_2 is *max*

Only x_1, x_3, x_5, x_6 may be the second largest key.

Analysis of Finding the 2nd

- Any algorithm that finds *secondLargest* must also find *max* before.   (*n*-1)

- The *secondLargest* can only be in those which lose directly to *max*.

- On its path along which bubbling up to the root of tournament tree, *max* beat $\lceil \lg n \rceil$ keys at most.

- Pick up *secondLargest*   ($\lceil \log n \rceil$-1)

- Total cost: *n*+$\lceil \log n \rceil$-2

Lower Bound by Adversary

- **Theorem**

  o Any algorithm (that works by comparing keys) to find the second largest in a set of n keys must do at least $n + \lceil \log n \rceil - 2$ comparisons in the worst case.

- **Proof**

  o There is an adversary strategy that can force any algorithm that finds *secondLargest* to compare *max* to $\lceil \log n \rceil$ distinct keys.

Weighted Key

- **Assigning a weight $w(x)$ to each key**
  - The initial values are all 1.
- **Adversary strategy**

Note: for one comparison, the weight increasing is no more than doubled.

Case	Adversary reply	Updating of weights
$w(x)>w(y)$	$x>y$	$w(x):=w(x)+w(y)$; $w(y):=0$
$w(x)=w(y)>0$	$x>y$	$w(x):=w(x)+w(y)$; $w(y):=0$
$w(y)>w(x)$	$y>x$	$w(y):=w(x)+w(y)$; $w(x):=0$
$w(x)=w(y)=0$	Consistent with previous replies	No change

Zero=Loss

Lower Bound by Adversary: Details

- **Note: the sum of weights is always n.**
- **Let x is *max*, then x is the only nonzero weighted key, that is $w(x)=n$.**
- **By the adversary rules:**

$$w_k(x) \leq 2w_{k-1}(x)$$

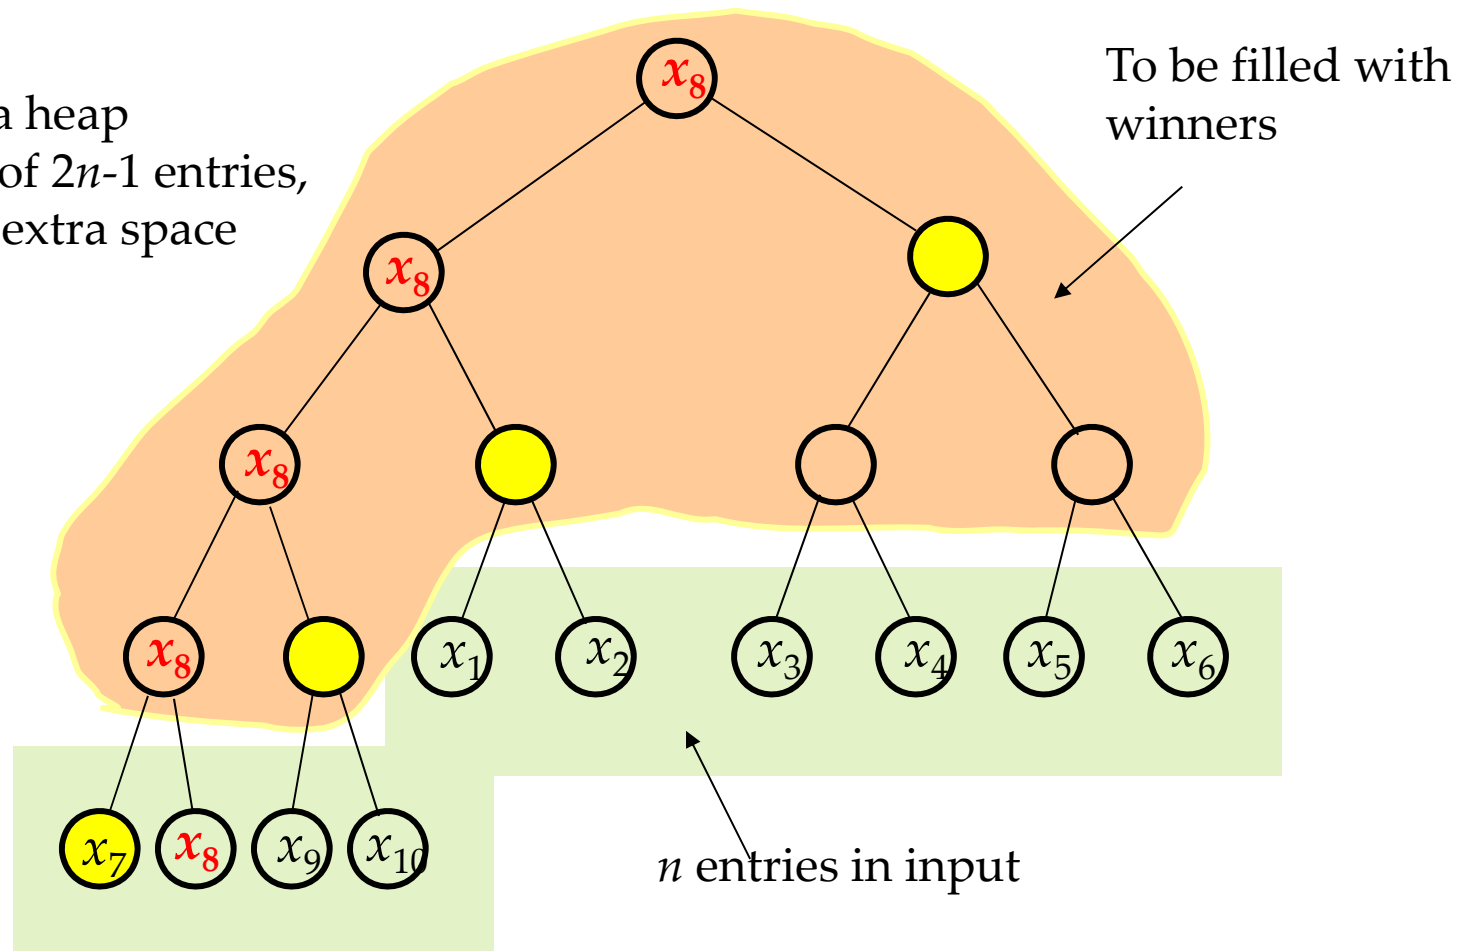- **Let K be the number of comparisons x wins against previously undefeated keys:**

$$n = w_K(x) \leq 2^K w_0(x) = 2^K$$

- **So, $K \geq \lceil \log n \rceil$**

Tracking the Losers to *MAX*

Building a heap
structure of $2n$-1 entries,
using n-1 extra space

To be filled with
winners

n entries in input

Finding the Median: the Strategy

- **Observation**
  - o If we can partition the problem set of keys into 2 subsets: S1, S2, such that any key in S1 is smaller that that of S2, the median must located in the set with more elements.

- **Divide-and-Conquer**
  - o Only one subset is needed to be processed recursively.

Adjusting the Rank

- **The rank of the median (of the original set) in the subset considered can be evaluated easily.**

- **An example**
  - Let n=255
  - The rank of median we want is 128
  - Assuming $|S_1|$=96, $|S_2|$=159
  - Then, the original median is in S_2, and the new rank is 128-96=32

Partitioning: Larger and Smaller

- **Dividing the array to be considered into two subsets: "small" and "large", the one with more elements will be processed recursively.**

[splitPoint]:   **pivot**

A "bad" pivot will give a very uneven partition!

small   large

for any element in this segment, the key is *less than pivot*.

To be   p rocessed   recursively

for any element in this segment, the key is *not less than pivot*.

Selection: the Algorithm

- **Input: S, a set of n keys; and k, an integer such that $1 \leq k \leq n$.**

- **Output: The kth smallest key in S.**

- **Note: Median selection is only a special case of the algorithm, with $k = \lceil n/2 \rceil$.**

- **Procedure**

- **Element select(SetOfElements S, int k)**
    - **if** $(|S| \leq 5)$ **return** direct solution; **else**
    - Constructing the subsets S_1 and $S_{2;}$
    - Processing one of S_1, S_2 with more elements, recursively.

Key issue:

How to construct the **partition** ?

Partition improved: the Strategy

greater than the median

less than the median of medians

All the elements are put in groups of 5

A

B
$>m^*$

m^*

C
$<m^*$

D

Medians

greater than the median

less than the median

Constructing the Partition

- **Find the m^*, the median of medians of all the groups of 5, as illustrated previously.**
- **Compare each key in sections A and D to m^*, and**
  - Let $S_1 = C \cup \{x \mid x \in A \cup D \text{ and } x < m^*\}$
  - Let $S_2 = B \cup \{x \mid x \in A \cup D \text{ and } x > m^*\}$

  (m^* is to be used as the pivot for the partition)

Divide and Conquer

if ($k=|S_1|+1$)

return m^*;

else if ($k \leq |S_1|$)

return select(S_1,k); //recursion

else

return select($S_2,k-|S_1|-1$); //recursion

Analysis

- **For simplicity:**

  o Assuming $n=5(2r+1)$ for all calls of *select*.

- $$W(n) \leq 6\left(\frac{n}{5}\right) + W\left(\frac{n}{5}\right) + 4r + W(7r+2)$$

The extreme case: all the elements in A∪D in one subset.
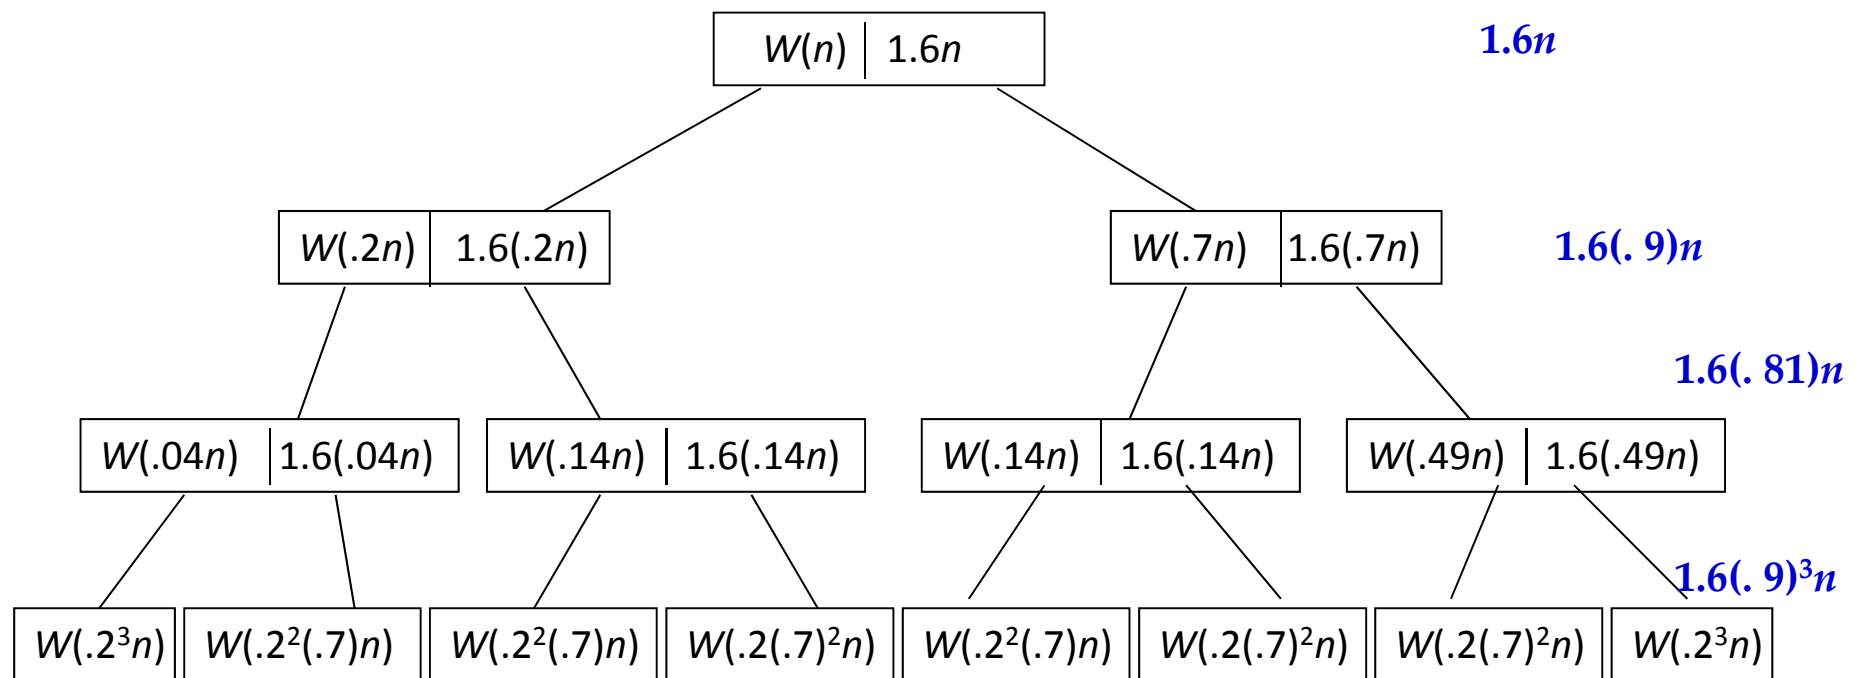
Finding the median in every group of 5

Finding the median of the medians

Comparing all the elements in A∪D with m^*

- ***Note: r is about n/10, and 0.7n+2 is about 0.7n***, **so**

$$W(n) \leq 1.6n + W(0.2n) + W(0.7n)$$

Worst Case Complexity of *Select*

$$W(n) \mid 1.6n$$

1.6n

$$W(.2n) \mid 1.6(.2n) \qquad W(.7n) \mid 1.6(.7n)$$

1.6(. 9)n

1.6(. 81)n

$$W(.04n) \mid 1.6(.04n) \qquad W(.14n) \mid 1.6(.14n) \qquad W(.14n) \mid 1.6(.14n) \qquad W(.49n) \mid 1.6(.49n)$$

1.6(. 9)³n

$$W(.2^3n) \quad W(.2^2(.7)n) \quad W(.2^2(.7)n) \quad W(.2(.7)^2n) \quad W(.2^2(.7)n) \quad W(.2(.7)^2n) \quad W(.2(.7)^2n) \quad W(.2^3n)$$

**Note: Row sums is a decreasing geometric series, so**
$$W(n) \in \Theta(n)$$

Relation to Median

- ## Observation
  - o Any algorithm of selection must know the relation of every element to the *median*.

Median

The adversary makes you wrong in either case

Crucial Comparison

- **A crucial comparison**
  - Establishing the relation of some x to the median.

- **Definition (for a comparison involving a key x)**
  - Crucial comparison for x: the first comparison where $x>y$, for some $y \geq$ median, or $x<y$ for some $y \leq$ median
  - Non-crucial comparison: the comparison between x and y where $x>$ median and $y<$ median, or vise versa

Adversary for Lower Bound

- **Status of the key during the running of the Algorithm:**
  - o *L*: Has been assigned a value *larger* than median
  - o *S*: Has been assigned a value *smaller* than median
  - o *N*: Has not yet been in a comparison

- **Adversary rule:**

Comparands	Adversary's action
N,N	one *L*, the another *S*
L,N or *N,L*	change *N* to *S*
S,N or *N,S*	change *N* to *L*

(In all other cases, just keep consistency)

Notes on the Adversary Arguments

- **All actions explicitly specified above make the comparisons un-crucial.**
  - At least, $(n\text{-}1)/2$ L or S can be assigned freely.
  - If there are already $(n\text{-}1)/2$ S, a value larger than median must be assigned to the new key, and if there are already $(n\text{-}1)/2$ L, a value smaller than median must be assigned to the new key. The last assigned value is the median.

- **So, an adversary can force the algorithm to do $(n\text{-}1)/2$ un-crucial comparisons at least(In the case that the algorithm start out by doing $(n\text{-}1)/2$ comparisons involving two N.**

Lower Bound for Selection Problem

- **Theorem:**
  - Any algorithm to find the median of n keys(for odd n) by comparison of keys must do at least 3n/2-3/2 comparisons in the worst case.

- **Argument:**
  - There must be done n-1 crucial comparisons at least.
  - An adversary can force the algorithm to perform as many as (*n*-1)/2 uncrucial comparisons.
    - Note: the algorithm can always start out by doing (*n*-1)/2 comparisons involving 2 *N*-keys, so, only (n-1)/2 *L* or *S* left for the adversary to assign freely as the adversary rule.

Thank you!

Q & A

Yu Huang

http://cs.nju.edu.cn/yuhuang