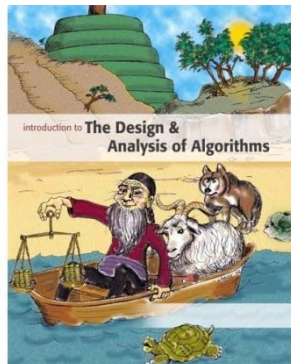




Introduction to

Algorithm Design and Analysis

[9] Hashing



Yu Huang

<http://cs.nju.edu.cn/yuhuang>
Institute of Computer Software
Nanjing University



In the Last Class...

- The searching problem
 - “Architecture” of data
- *log_n* search
 - Binary search
 - In a more general sense
 - Red-black tree: balanced BST
 - Definition
 - Black height constraint for balance
 - Color constraint for low maintenance cost
 - Operation
 - Insertion, deletion

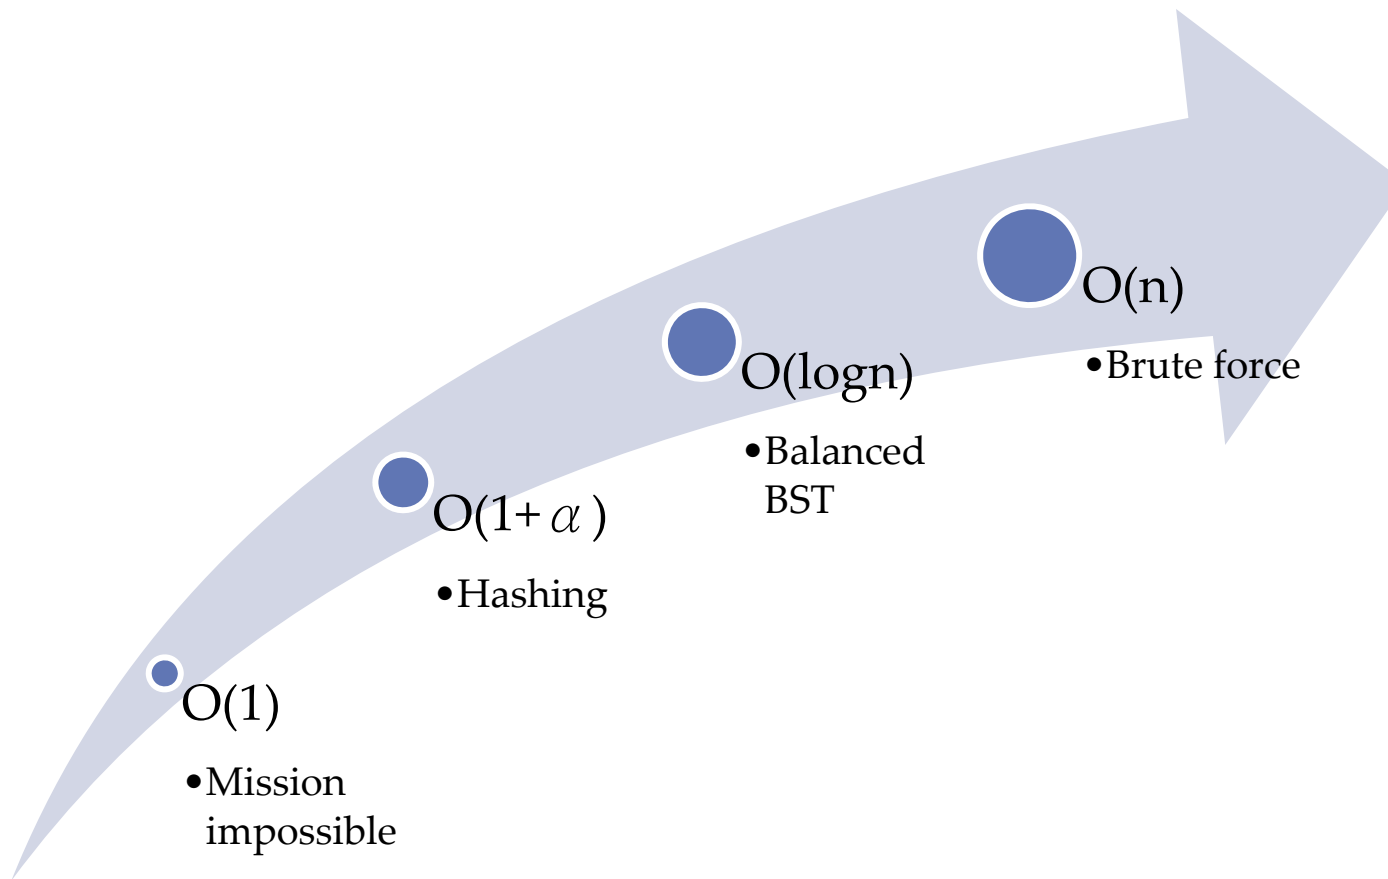


Hashing

- **The searching problem**
 - The ambition of hashing
- **Hashing**
 - Brute force table: direct addressing
 - Basic idea of hashing
- **Collision Handling for Hashing**
 - Closed address hashing
 - Open address hashing
- **Amortized Analysis**
 - Array doubling



Cost for Searching




Cost for Searching

- **Brute force**
 - $O(n)$
- **Balanced BST**
 - $O(\log n)$
- **Hashing – almost constant time**
 - $O(1+\alpha)$
- **“Mission impossible”**
 - $O(1)$



Searching - a Brute Force Approach

- Direct-address table
 - Take into account the *whole universe* of keys



Direct-address Table


DIRECT-ADDRESS-SEARCH (T, k)
return $T[k]$

DIRECT-ADDRESS-INSERT (T, x)
 $T[\text{key}[x]] := x$


DIRECT-ADDRESS-DELETE (T, x)
 $T[\text{key}[x]] := \text{NIL}$

Hashing: the Idea

Hash Table (in feasible size)




- **Index distribution**
- **Collision handling**



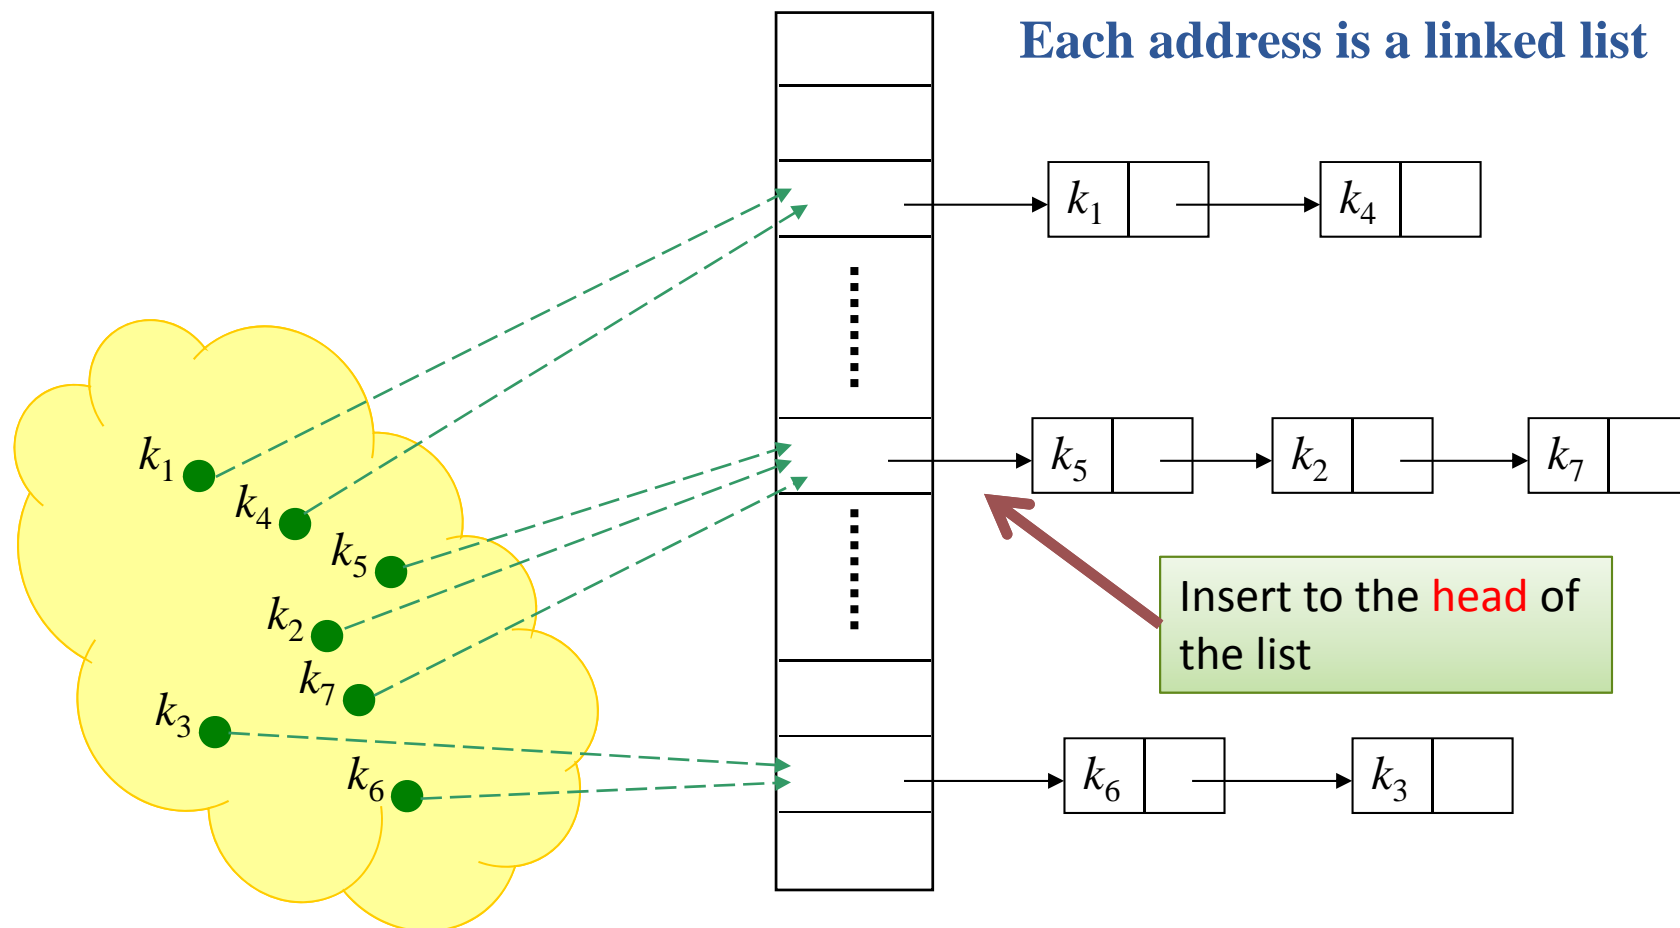
A calculated
array index for
the key

Very large, but only a
small part is used in an
application



Value of a
specific key

Collision Handling: Closed Address



Closed Address - Analysis

- **Assumption - simple uniform hashing**
 - For $j=0,1,2,\dots,m-1$, the average length of the list at $E[j]$ is n/m .
- **The average cost for an unsuccessful search**
 - Any key that is not in the table is equally likely to hash to any of the m address.
 - Total cost $\Theta(1 + n/m)$
 - The average cost to determine that the key is not in the list $E[h(k)]$ is the cost to search to the end of the list, which is n/m .



Closed Address - Analysis

- **For successful search** (assuming that x_i is the i^{th} element inserted into the table, $i=1,2,\dots,n$)
 - For each i , the probability of that x_i is searched is $1/n$.
 - For a specific x_i , the number of elements examined in a successful search is $t+1$, where t is the number of elements inserted into the same list as x_i , after x_i has been inserted

$$\frac{1}{n} \sum_{i=1}^n (1 + t)$$

- **How to compute t ?**
 - Consider the *construction* process of the hash table



Closed Address - Analysis

- **For successful search:** (assuming that x_i is the i^{th} element inserted into the table, $i=1,2,\dots,n$)
 - For each i , the probability of that x_i is searched is $1/n$.
 - For a specific x_i , the number of elements examined in a successful search is $t+1$, where t is the number of elements inserted into the same list as x_i , after x_i has been inserted. And for any j , the probability of that x_j is inserted into the same list of x_i is $1/m$. So, the cost is:

Cost for
computing
hashing

$$1 + \frac{1}{n} \sum_{i=1}^n \left(1 + \sum_{j=i+1}^n \frac{1}{m} \right)$$

Expected number of
elements in front of the
searched one in the same
linked list.



Closed Address: Analysis

- The average cost of a successful search:
 - Define $\alpha = n/m$ as *load factor*,

The average cost of a successful search is :

$$\frac{1}{n} \sum_{i=1}^n \left(1 + \sum_{j=i+1}^n \frac{1}{m} \right) = 1 + \frac{1}{nm} \sum_{i=1}^n (n-i) = 1 + \frac{1}{nm} \sum_{i=1}^{n-1} i$$
$$= 1 + \frac{n-1}{2m} = 1 + \frac{\alpha}{2} - \frac{\alpha}{2n} = \Theta(1 + \alpha)$$

Number of elements in front of the searched one in the same linked list

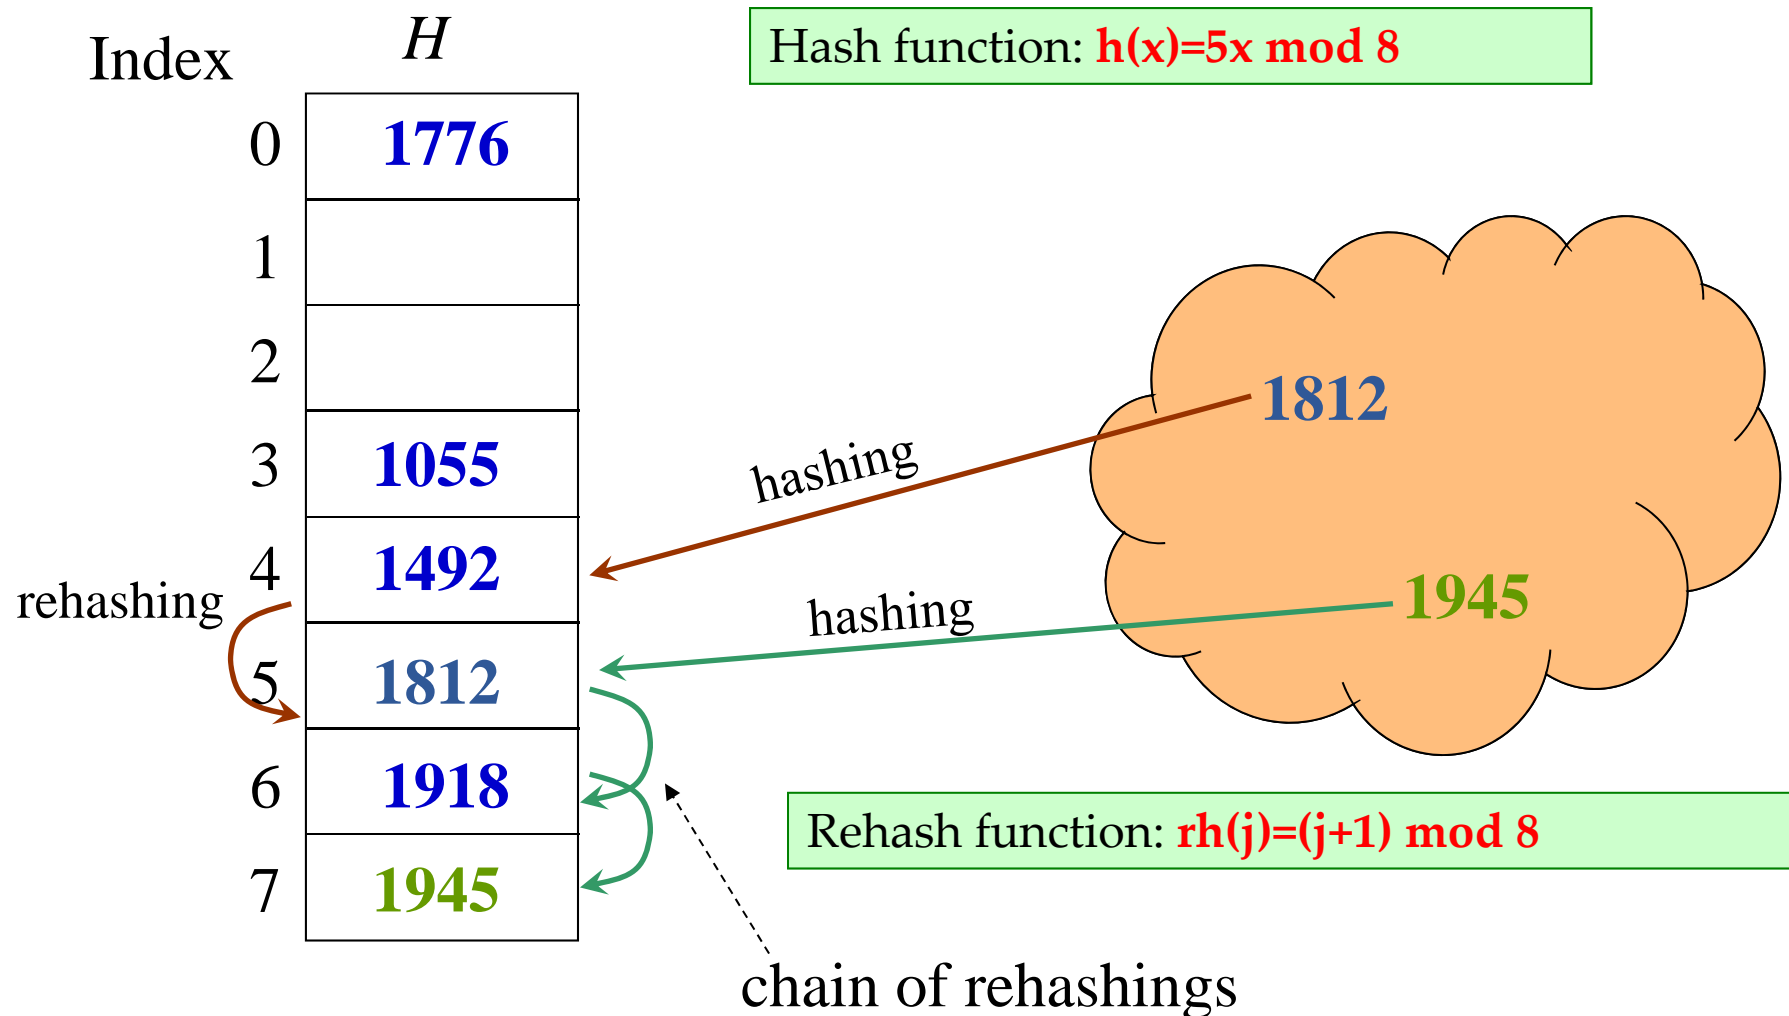


Collision Handling: Open Address

- All elements are stored in the hash table
 - No linked list is used
 - The load factor α cannot be larger than 1
- Collision is settled by “rehashing”
 - A function is used to get a new hashing address for each collided address
 - The hash table slots are *probed* successively, until a valid location is found.
- The probe sequence can be seen as a permutation of $(0, 1, 2, \dots, m-1)$



Linear Probing: An Example



Commonly Used Probing

Linear probing:

Given an ordinary hash function h' , which is called an auxiliary hash function, the hash function is: (clustering may occur)

$$h(k,i) = (h'(k) + i) \bmod m \quad (i=0,1,\dots,m-1)$$

Quadratic Probing:

Given auxiliary function h' and nonzero auxiliary constant c_1 and c_2 , the hash function is: (secondary clustering may occur)

$$h(k,i) = (h'(k) + c_1 i + c_2 i^2) \bmod m \quad (i=0,1,\dots,m-1)$$

Double hashing:

Given auxiliary functions h_1 and h_2 , the hash function is:

$$h(k,i) = (h_1(k) + i h_2(k)) \bmod m \quad (i=0,1,\dots,m-1)$$



Equally Likely Permutations

- **Assumption**

- Each key is equally likely to have any of the $m!$ permutations of $(1, 2, \dots, m)$ as its probe sequence

- **Note**

- Both linear and quadratic probing have only m distinct probe sequence, as determined by the first probe



Analysis for Open Address Hashing

- The average number of probes in an unsuccessful search is at most $1/(1-\alpha)$ ($\alpha=n/m<1$)
 - Assuming uniform hashing

The probability of the first probed position being occupied is $\frac{n}{m}$, and that of the j^{th} ($j > 1$) position occupied is $\frac{n-j+1}{m-j+1}$. So the probability of the number of probes no less than i will be:

$$\frac{n}{m} \cdot \frac{n-1}{m-1} \cdot \frac{n-2}{m-2} \cdots \frac{n-i+2}{m-i+2} \leq \left(\frac{n}{m}\right)^{i-1} = \alpha^{i-1}$$

The the average number of probe is: $\sum_{i=1}^{\infty} \alpha^{i-1} = \sum_{i=0}^{\infty} \alpha^i = \frac{1}{1-\alpha}$

See [CLRS] p.1199, C.25



Analysis for Open Address Hashing

- The average cost of probes in an successful search is at most $\frac{1}{\alpha} \ln \frac{1}{1-\alpha}$ ($\alpha=n/m<1$)

- Assuming uniform hashing

To search for the $(i+1)^{th}$ inserted element in the table, the cost is the same as that for inserting it when there are just i elements in the table.

At that time, $\alpha = \frac{i}{m}$. So the cost is $\frac{1}{1-\frac{i}{m}} = \frac{m}{m-i}$.

So the average cost for a successful search is:

$$\begin{aligned} \frac{1}{n} \sum_{i=0}^{n-1} \frac{m}{m-i} &= \frac{m}{n} \sum_{i=0}^{n-1} \frac{1}{m-i} = \frac{1}{\alpha} \sum_{i=m-n+1}^m \frac{1}{i} \\ &\leq \frac{1}{\alpha} \int_{m-n}^m \frac{dx}{x} = \frac{1}{\alpha} \ln \frac{m}{m-n} = \frac{1}{\alpha} \ln \frac{1}{1-\alpha} \end{aligned}$$

For your reference:

Half full: 1.387;

90% full: 2.559



Hash Function

- **A good hash function satisfies the assumption of simple uniform hashing**
 - Heuristic hashing functions
 - The division method: $h(k) = k \bmod m$
 - The multiplication method: $h(k) = \lfloor m(kA \bmod 1) \rfloor$ ($0 < A < 1$)
 - No single function can avoid the worst case $\Theta(n)$.
 - So “universal hashing” is proposed.
 - Rich resource about hashing function
 - Gonnet and Baeza-Yates: Handbook of Algorithms and Data Structures, Addison-Wesley, 1991.



Array Doubling

- Cost for search in a hash table is $\Theta(1+\alpha)$
 - If we can keep α constant, the cost will be $\Theta(1)$
- What if the hash table is more and more loaded?
 - Space allocation techniques such as array doubling may be needed
- The problem of “**unusually expensive**” individual operation



Looking at the Memory Allocation

- `hashingInsert(HASHTABLE H , ITEM x)`
 - `integer $size=0$, $num=0$;`
 - `if $size=0$ then allocate a block of size 1; $size=1$;`
 - `if $num=size$ then`
 - `allocate a block of size $2size$;`
 - `move all item into new table;`
 - `$size=2size$;`
 - `insert x into the table;`
 - `$num=num+1$;`
 - `return`
- Insertion with expansion: cost $size$
- Elementary insertion: cost 1

Worst-case Analysis

- For n execution of insertion operations
 - A bad analysis: the worst case for one insertion is the case when expansion is required, up to n
 - So, the worst case cost is in $O(n^2)$.
- Note the expansion is required during the i th operation only if $i=2^k$, and the cost of the i th operation

$$c_i = \begin{cases} i & \text{if } i - 1 \text{ is exactly the power of } 2 \\ 1 & \text{otherwise} \end{cases}$$

So the total cost is: $\sum_{i=1}^n c_i \leq n + \sum_{j=0}^{\lfloor \log n \rfloor} 2^j < n + 2n = 3n$



Amortized Analysis – Why?

- Unusually expensive operations
 - E.g., Insert-with-array-doubling
- **Relation** between expensive and usual operations
 - Each piece of the doubling cost corresponds to some previous insert




Amortized Analysis

- **Amortized equation:**
amortized cost = actual cost + accounting cost
- **Design goals for accounting cost**
 - In **any** legal sequence of operations, the sum of the accounting costs is nonnegative
 - The amortized cost of each operation is fairly regular, in spite of the wide fluctuate possible for the actual cost of individual operations



Amortized Analysis: MultiPop Stack



Amortized cost: push:2; pop, multipop: 0



Amortized Analysis: Binary Counter

0	000000000	0
1	000000001	1
2	000000010	3
3	000000011	4
4	000000100	7
5	000000101	8
6	000000110	10
7	000000111	11
8	000001000	15
9	000001001	16
10	000001010	18
11	000001011	19
12	000001100	22
13	000001101	23
14	000001110	25
15	000001111	26
16	000100000	31

Cost measure: bit flip

amortized cost:

set 1: 2

set 0: 0



Accounting Scheme for Stack Push

- Push operation with array doubling
 - No resize triggered: 1
 - Resize($n \rightarrow 2n$) triggered: $nt+1$ (t is a constant)
- Accounting scheme (specifying accounting cost)
 - No resize triggered: $2t$
 - Resize($n \rightarrow 2n$) triggered: $-nt+2t$
- So, the amortized cost of each individual push operation is $1+2t \in \Theta(1)$



Thank you!

Q & A

Yu Huang

<http://cs.nju.edu.cn/yuhuang>

